

Airflow Installation

To install airflow we need docker-compose file that's available in official airflow documentation.

Follow this steps to install configure and test airflow

<https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html>

1. Create a folder called "airflow"
2. Past docker compose file
3. Docker compose up -d
4. Airflow should run on port 8080
5. Check all the volume and port details on the docker compose file

Testing airflow

1. Open airflow ui using port 8080
2. User & Pass: airflow
3. Check the terminal (configs dags logs plugins) - folders are available
4. Create an simple dag on the dags folder (scheduled task)
5. Verify the dags are visible on the ui - and run it - check logs

Docker compose file

docker-compose.yaml

```
# Licensed to the Apache Software Foundation (ASF) under one
# or more contributor license agreements. See the NOTICE file
# distributed with this work for additional information
# regarding copyright ownership. The ASF licenses this file
# to you under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance
```

```
# with the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing,
# software distributed under the License is distributed on an
# "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
# KIND, either express or implied. See the License for the
# specific language governing permissions and limitations
# under the License.
#

# Basic Airflow cluster configuration for CeleryExecutor with Redis and PostgreSQL.
#
# WARNING: This configuration is for local development. Do not use it in a production
# deployment.
#
# This configuration supports basic configuration using environment variables or an .env
# file
# The following variables are supported:
#
# AIRFLOW_IMAGE_NAME          - Docker image name used to run Airflow.
#                               Default: apache/airflow:3.0.4
# AIRFLOW_UID                 - User ID in Airflow containers
#                               Default: 50000
# AIRFLOW_PROJ_DIR            - Base path to which all the files will be volumed.
#                               Default: .
# Those configurations are useful mostly in case of standalone testing/running Airflow in
# test/try-out mode
#
# _AIRFLOW_WWW_USER_USERNAME - Username for the administrator account
# (if requested).
#                               Default: airflow
# _AIRFLOW_WWW_USER_PASSWORD - Password for the administrator account (if
# requested).
#                               Default: airflow
# _PIP_ADDITIONAL_REQUIREMENTS - Additional PIP requirements to add when
# starting all containers.
#                               Use this option ONLY for quick checks. Installing requirements at
# container
```

```

#           startup is done EVERY TIME the service is started.
#           A better way is to build a custom image or extend the official
image
#           as described in
https://airflow.apache.org/docs/docker-stack/build.html.
#           Default: "
#
# Feel free to modify this file to suit your needs.
---
x-airflow-common:
  &airflow-common
  # In order to add custom dependencies or upgrade provider distributions you can use
  your extended image.
  # Comment the image line, place your Dockerfile in the directory where you placed the
  docker-compose.yaml
  # and uncomment the "build" line below, Then run `docker-compose build` to build the
  images.
  image: ${AIRFLOW_IMAGE_NAME:-apache/airflow:3.0.4}
  # build: .
  environment:
    &airflow-common-env
    AIRFLOW__CORE__EXECUTOR: CeleryExecutor
    AIRFLOW__CORE__AUTH_MANAGER:
airflow.providers.fab.auth_manager.fab_auth_manager.FabAuthManager
    AIRFLOW__DATABASE__SQL_ALCHEMY_CONN:
postgres+psycopg2://airflow:airflow@postgres/airflow
    AIRFLOW__CELERY__RESULT_BACKEND:
db+postgres://airflow:airflow@postgres/airflow
    AIRFLOW__CELERY__BROKER_URL: redis://:@redis:6379/0
    AIRFLOW__CORE__FERNET_KEY: "
    AIRFLOW__CORE__DAGS_ARE_PAUSED_AT_CREATION: 'true'
    AIRFLOW__CORE__LOAD_EXAMPLES: 'true'
    AIRFLOW__CORE__EXECUTION_API_SERVER_URL:
'http://airflow-apiserver:8080/execution/'
    # yamllint disable rule:line-length
    # Use simple http server on scheduler for health checks
    # See
https://airflow.apache.org/docs/apache-airflow/stable/administration-and-deployment/log
ging-monitoring/check-health.html#scheduler-health-check-server
    # yamllint enable rule:line-length

```

```

AIRFLOW__SCHEDULER__ENABLE_HEALTH_CHECK: 'true'
# WARNING: Use _PIP_ADDITIONAL_REQUIREMENTS option ONLY for a quick
checks
# for other purpose (development, test and especially production usage) build/extend
Airflow image.
_PIP_ADDITIONAL_REQUIREMENTS: ${_PIP_ADDITIONAL_REQUIREMENTS:-}
# The following line can be used to set a custom config file, stored in the local config
folder
AIRFLOW_CONFIG: '/opt/airflow/config/airflow.cfg'
volumes:
- ${AIRFLOW_PROJ_DIR:-.}/dags:/opt/airflow/dags
- ${AIRFLOW_PROJ_DIR:-.}/logs:/opt/airflow/logs
- ${AIRFLOW_PROJ_DIR:-.}/config:/opt/airflow/config
- ${AIRFLOW_PROJ_DIR:-.}/plugins:/opt/airflow/plugins
user: "${AIRFLOW_UID:-50000}:0"
depends_on:
&airflow-common-depends-on
redis:
condition: service_healthy
postgres:
condition: service_healthy

services:
postgres:
image: postgres:13
environment:
POSTGRES_USER: airflow
POSTGRES_PASSWORD: airflow
POSTGRES_DB: airflow
volumes:
- postgres-db-volume:/var/lib/postgresql/data
healthcheck:
test: ["CMD", "pg_isready", "-U", "airflow"]
interval: 10s
retries: 5
start_period: 5s
restart: always

redis:
# Redis is limited to 7.2-bookworm due to licencing change

```

<https://redis.io/blog/redis-adopts-dual-source-available-licensing/>

image: redis:7.2-bookworm

expose:

- 6379

healthcheck:

test: ["CMD", "redis-cli", "ping"]

interval: 10s

timeout: 30s

retries: 50

start_period: 30s

restart: always

airflow-apiserver:

<<: *airflow-common

command: api-server

ports:

- "8080:8080"

healthcheck:

test: ["CMD", "curl", "--fail", "http://localhost:8080/api/v2/version"]

interval: 30s

timeout: 10s

retries: 5

start_period: 30s

restart: always

depends_on:

<<: *airflow-common-depends-on

airflow-init:

condition: service_completed_successfully

airflow-scheduler:

<<: *airflow-common

command: scheduler

healthcheck:

test: ["CMD", "curl", "--fail", "http://localhost:8974/health"]

interval: 30s

timeout: 10s

retries: 5

start_period: 30s

restart: always

depends_on:

```
<<: *airflow-common-depends-on
airflow-init:
  condition: service_completed_successfully
```

```
airflow-dag-processor:
  <<: *airflow-common
  command: dag-processor
  healthcheck:
    test: ["CMD-SHELL", 'airflow jobs check --job-type DagProcessorJob --hostname
    "${HOSTNAME}"]
    interval: 30s
    timeout: 10s
    retries: 5
    start_period: 30s
  restart: always
  depends_on:
    <<: *airflow-common-depends-on
  airflow-init:
    condition: service_completed_successfully
```

```
airflow-worker:
  <<: *airflow-common
  command: celery worker
  healthcheck:
    # yamllint disable rule:line-length
    test:
      - "CMD-SHELL"
      - 'celery --app airflow.providers.celery.executors.celery_executor.app inspect ping
      -d "celery@${HOSTNAME}" || celery --app airflow.executors.celery_executor.app
      inspect ping -d "celery@${HOSTNAME}"'
    interval: 30s
    timeout: 10s
    retries: 5
    start_period: 30s
  environment:
    <<: *airflow-common-env
    # Required to handle warm shutdown of the celery workers properly
    # See
    https://airflow.apache.org/docs/docker-stack/entrypoint.html#signal-propagation
    DUMB_INIT_SETSID: "0"
```

```
restart: always
depends_on:
  <<: *airflow-common-depends-on
airflow-apiserver:
  condition: service_healthy
airflow-init:
  condition: service_completed_successfully
```

```
airflow-triggerer:
  <<: *airflow-common
command: triggerer
healthcheck:
  test: ["CMD-SHELL", 'airflow jobs check --job-type TriggererJob --hostname
"${HOSTNAME}"]
  interval: 30s
  timeout: 10s
  retries: 5
  start_period: 30s
restart: always
depends_on:
  <<: *airflow-common-depends-on
airflow-init:
  condition: service_completed_successfully
```

```
airflow-init:
  <<: *airflow-common
entrypoint: /bin/bash
# yamllint disable rule:line-length
command:
  - -c
  - |
    if [[ -z "${AIRFLOW_UID}" ]]; then
      echo
      echo -e "\033[1;33mWARNING!!!: AIRFLOW_UID not set!\e[0m"
      echo "If you are on Linux, you SHOULD follow the instructions below to set "
      echo "AIRFLOW_UID environment variable, otherwise files will be owned by
root."
      echo "For other operating systems you can get rid of the warning with manually
created .env file:"
```

```

    echo "    See:
https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html
#setting-the-right-airflow-user"
    echo
    export AIRFLOW_UID=$(id -u)
fi
one_meg=1048576
mem_available=$((($(getconf _PHYS_PAGES) * $(getconf PAGE_SIZE) /
one_meg))
cpus_available=$(grep -cE 'cpu[0-9]+' /proc/stat)
disk_available=$((df / | tail -1 | awk '{print $4}'))
warning_resources="false"
if (( mem_available < 4000 )) ; then
    echo
    echo -e "\033[1;33mWARNING!!!: Not enough memory available for
Docker.\e[0m"
    echo "At least 4GB of memory required. You have $(numfmt --to iec
$$(mem_available * one_meg)))"
    echo
    warning_resources="true"
fi
if (( cpus_available < 2 )); then
    echo
    echo -e "\033[1;33mWARNING!!!: Not enough CPUS available for Docker.\e[0m"
    echo "At least 2 CPUs recommended. You have ${cpus_available}"
    echo
    warning_resources="true"
fi
if (( disk_available < one_meg * 10 )); then
    echo
    echo -e "\033[1;33mWARNING!!!: Not enough Disk space available for
Docker.\e[0m"
    echo "At least 10 GBs recommended. You have $(numfmt --to iec
$$(disk_available * 1024 ))"
    echo
    warning_resources="true"
fi
if [[ ${warning_resources} == "true" ]]; then
    echo

```



```
    echo -e "\033[1;33mWARNING!!!: You have not enough resources to run Airflow
(see above)\n\e[0m"
    echo "Please follow the instructions to increase amount of resources available:"
    echo "
https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html
#before-you-begin"
    echo
fi
echo
echo "Creating missing opt dirs if missing:"
echo
mkdir -v -p /opt/airflow/{logs,dags,plugins,config}
echo
echo "Airflow version:"
/entrypoint airflow version
echo
echo "Files in shared volumes:"
echo
ls -la /opt/airflow/{logs,dags,plugins,config}
echo
echo "Running airflow config list to create default config file if missing."
echo
/entrypoint airflow config list >/dev/null
echo
echo "Files in shared volumes:"
echo
ls -la /opt/airflow/{logs,dags,plugins,config}
echo
echo "Change ownership of files in /opt/airflow to ${AIRFLOW_UID}:0"
echo
chown -R "${AIRFLOW_UID}:0" /opt/airflow/
echo
echo "Change ownership of files in shared volumes to ${AIRFLOW_UID}:0"
echo
chown -v -R "${AIRFLOW_UID}:0" /opt/airflow/{logs,dags,plugins,config}
echo
echo "Files in shared volumes:"
echo
ls -la /opt/airflow/{logs,dags,plugins,config}
```

```
# yamllint enable rule:line-length
environment:
  <<: *airflow-common-env
  _AIRFLOW_DB_MIGRATE: 'true'
  _AIRFLOW_WWW_USER_CREATE: 'true'
  _AIRFLOW_WWW_USER_USERNAME:
${_AIRFLOW_WWW_USER_USERNAME:-airflow}
  _AIRFLOW_WWW_USER_PASSWORD:
${_AIRFLOW_WWW_USER_PASSWORD:-airflow}
  _PIP_ADDITIONAL_REQUIREMENTS: "
  user: "0:0"
```

```
airflow-cli:
  <<: *airflow-common
  profiles:
    - debug
  environment:
    <<: *airflow-common-env
    CONNECTION_CHECK_MAX_COUNT: "0"
  # Workaround for entrypoint issue. See:
https://github.com/apache/airflow/issues/16252
  command:
    - bash
    - -c
    - airflow
  depends_on:
    <<: *airflow-common-depends-on
```

You can enable flower by adding "--profile flower" option e.g. docker-compose
--profile flower up

or by explicitly targeted on the command line e.g. docker-compose up flower.

See: <https://docs.docker.com/compose/profiles/>

flower:

```
<<: *airflow-common
command: celery flower
profiles:
  - flower
ports:
  - "5555:5555"
healthcheck:
```

```
test: ["CMD", "curl", "--fail", "http://localhost:5555/"]
interval: 30s
timeout: 10s
retries: 5
start_period: 30s
restart: always
depends_on:
  <<: *airflow-common-depends-on
airflow-init:
  condition: service_completed_successfully
```

```
volumes:
  postgres-db-volume:
```

Docker compose up -d