

Kafka

To install Kafk we need a docker-compose file (compose file created by chat gpt).
In compose file installed zookeeper, kafka, kafka-ui

Follow this steps to install configure and test airflow

1. Create a folder called "kafka"
2. Past docker compose file script
3. Docker compose up -d
4. Kafka ui should run on port 8082
5. Check all the volume and port details on the docker compose file

Testing kafka

1. Open airflow ui using port 8082
2. Try to communicate send and receive to verify
3. Check in kafka-ui the topics and messages are there

Installation Steps

1. **Create project folder & file**

Shell

```
mkdir kafka-cluster  
cd kafka-cluster  
nano docker-compose.yml
```

2. **Paste the above YAML and save.**

3. **Start the cluster**

Shell

```
docker compose up -d
```

4. Verify services

Shell

```
docker ps
```

You should see `zookeeper` and `kafka` running.

5. Check logs

Shell

```
docker logs kafka | grep -i started
```

Look for `Kafka Server started`.

Test Kafka Setup

Create a topic:

Shell

```
docker exec -it kafka bash
kafka-topics --create --topic test-topic --bootstrap-server
localhost:9092 --partitions 1 --replication-factor 1
```

List topics:

Shell

```
kafka-topics --list --bootstrap-server localhost:9092
```

Produce messages:

Shell

```
kafka-console-producer --topic test-topic  
--bootstrap-server localhost:9092
```

(Type a few lines, press **Enter** after each)

Consume messages:

Shell

```
kafka-console-consumer --topic test-topic  
--bootstrap-server localhost:9092 --from-beginning
```

Check Data Persistence

- Stop the cluster:

Shell

```
docker compose down
```

- Start again:

Shell

```
docker compose up -d
```

Your topics and data remain because of mounted volumes.

docker-compose.yml

services:

zookeeper:

image: confluentinc/cp-zookeeper:7.6.0
container_name: zookeeper
restart: unless-stopped
ports:
- "2181:2181"
environment:
 ZOOKEEPER_CLIENT_PORT: 2181
 ZOOKEEPER_TICK_TIME: 2000
volumes:
- zookeeper-data:/var/lib/zookeeper/data
- zookeeper-log:/var/lib/zookeeper/log
networks:
- kafka-net
healthcheck:
 test: ["CMD-SHELL", "bash -c '</dev/tcp/localhost/2181'"]
 interval: 10s
 timeout: 5s
 retries: 5

kafka:

image: confluentinc/cp-kafka:7.6.0
container_name: kafka
restart: unless-stopped
depends_on:
 zookeeper:
 condition: service_healthy
ports:
- "9092:9092" # Internal listener
- "29092:29092" # External listener
environment:
 KAFKA_BROKER_ID: 1
 KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
 KAFKA_LISTENERS:
PLAINTEXT://0.0.0.0:9092,PLAINTEXT_HOST://0.0.0.0:29092
 KAFKA_ADVERTISED_LISTENERS:
PLAINTEXT://kafka:9092,PLAINTEXT_HOST://localhost:29092
 KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
 KAFKA_INTER_BROKER_LISTENER_NAME: PLAINTEXT
 KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1

KAFKA_GROUP_INITIAL_REBALANCE_DELAY_MS: 0

volumes:

- kafka-data:/var/lib/kafka/data

networks:

- kafka-net

healthcheck:

test: ["CMD", "kafka-broker-api-versions", "--bootstrap-server=localhost:9092"]

interval: 15s

timeout: 10s

retries: 5

kafka-ui:

image: provectuslabs/kafka-ui:latest

container_name: kafka-ui

restart: unless-stopped

depends_on:

kafka:

condition: service_healthy

ports:

- "8082:8080" # Access UI on http://localhost:8082

environment:

KAFKA_CLUSTERS_0_NAME: local

KAFKA_CLUSTERS_0_BOOTSTRAPSERVERS: kafka:9092

KAFKA_CLUSTERS_0_ZOOKEEPER: zookeeper:2181

networks:

- kafka-net

healthcheck:

test: ["CMD", "curl", "-f", "http://localhost:8082"]

interval: 10s

timeout: 5s

retries: 5

volumes:

zookeeper-data:

zookeeper-log:

kafka-data:

networks:

kafka-net:

driver: bridge