

Top-level Linux Directories and Their Purpose

♦ /

- The **root directory** — everything starts here.
 - All other files and directories exist under /.
-

♦ /bin

- **Essential user binaries (commands)** required for all users.
 - Examples: `ls`, `cp`, `mv`, `cat`, `bash`.
 - Needed even if no other filesystems are mounted.
-

♦ /sbin

- **System binaries** (mostly for system administration).
 - Examples: `ifconfig`, `mount`, `shutdown`.
 - Only root/admin usually runs these.
-

♦ /etc

- **System configuration files** (text-based).
 - Example: `/etc/passwd` (user accounts), `/etc/ssh/sshd_config` (SSH config), `/etc/fstab` (mount info).
 - Think of `/etc` as the **settings folder** of the OS.
-

♦ `/var`

- **Variable data** that changes frequently.
 - Examples:
 - `/var/log` → system & app logs
 - `/var/spool` → print/mail queues
 - `/var/tmp` → temporary files that survive reboot
-

♦ `/lib` and `/lib64`

- **Shared libraries** needed by programs in `/bin` and `/sbin`.
 - Similar to DLLs on Windows.
 - Example: `/lib/x86_64-linux-gnu/libc.so.6`.
-

♦ `/usr`

- **User applications and utilities** (not essential for boot, but for normal usage).
 - Contains programs, libraries, docs.
 - Inside `/usr`:
 - `/usr/bin` → extra commands for users (`vim`, `python`, `git`)
 - `/usr/sbin` → extra system admin commands (`apache2`, `nginx`)
 - `/usr/lib` → libraries for above binaries
 - `/usr/share` → shared files (icons, docs, man pages)
-

♦ `/opt`

- **Optional software** (usually third-party apps).
 - Example: if you install Google Chrome manually, it may go in `/opt/google/chrome`.
-

♦ `/home`

- **User home directories.**
 - Example: `/home/vikram` for your files, downloads, configs.
 - Like "C:\Users" in Windows.
-

♦ `/root`

- The **home directory of the root user**.
 - Do not confuse with `/`.
-

♦ `/tmp`

- **Temporary files**, deleted after reboot.
 - All users and apps use it for scratch space.
-

♦ `/boot`

- Files needed for booting Linux.
 - Example: Kernel (`vmlinux`), initramfs, bootloader (GRUB config).
-

♦ `/dev`

- **Device files** (special files that represent hardware).
 - Example: `/dev/sda` (disk), `/dev/tty` (terminal), `/dev/null`.
-

♦ `/mnt`

- **Temporary mount point** for mounting filesystems (manual usage).
 - Example: If you mount a USB drive manually, you might use `/mnt/usb`.
-

♦ `/media`

- **Automatic mount point** for removable devices.
 - Example: Plugging in a USB drive → `/media/vikram/USB`.
-

♦ `/proc`

- **Virtual filesystem** that provides info about running processes and kernel.
 - Example: `/proc/cpuinfo`, `/proc/meminfo`, `/proc/1234/` (process with PID 1234).
 - It's not real files, just runtime info.
-

♦ `/sys`

- Another **virtual filesystem** to interact with the kernel and devices.
 - Example: `/sys/class/net/` shows network interfaces.
-

♦ `/srv`

- **Service data** served by the system.
 - Example: web server files may be stored in `/srv/www`.
-

♦ `/run`

- Stores **runtime process data** since last boot.
 - Example: PID files (`/run/nginx.pid`).
 - Cleared on reboot.
-

Summary (Easy Mapping)

- `/bin` → Basic commands
- `/sbin` → System commands
- `/etc` → Configs
- `/var` → Logs, caches, variable files
- `/lib` → Libraries
- `/usr` → Extra software
- `/opt` → Third-party apps
- `/home` → User files
- `/root` → Root user's home
- `/tmp` → Temporary files

- `/boot` → Boot loader & kernel
 - `/dev` → Devices
 - `/mnt`, `/media` → Mount points
 - `/proc`, `/sys` → Kernel & process info
 - `/srv` → Service data
 - `/run` → Runtime info
-

👉 Now, in companies (MNCs), when they deploy apps, they usually:

- Put **config** files in `/etc/appname/`
 - Put **binaries** in `/usr/bin` or `/opt/appname/`
 - Put **logs** in `/var/log/appname/`
 - Keep **data** in `/var/lib/appname/`
-

Best Way to Install 3rd Party Software with Docker Compose

1. Use `/srv` for all deployments

- Why: `/srv` is designed for **service data & deployments**.
- Each app gets its **own folder** inside `/srv`.
- Inside that folder, you keep:
 - `docker-compose.yml`
 - configs
 - persistent volumes

📁 Example layout:

```
None
/srv/
├── caddy/
│   ├── docker-compose.yml
│   ├── configs/           # app configs (e.g., Caddyfile)
│   ├── site/              # static site files
│   └── data/               # persistent container state
├── airflow/
│   ├── docker-compose.yml
│   ├── configs/           # airflow.cfg, env vars
│   └── logs/               # persistent logs
├── mlflow/
│   ├── docker-compose.yml
│   ├── configs/           # mlflow.conf, creds
│   └── data/               # experiments, models
```

2. Use named volumes for container internals

- Inside Compose, map container paths to **named volumes** → durability.
- Use **bind mounts only for configs & site files** (things you want to edit easily).

♦ Example (Caddy):

```
None
version: '3.9'
services:
  caddy:
    image: caddy:latest
    container_name: caddy
    restart: unless-stopped
    ports:
      - "80:80"
      - "443:443"
    volumes:
      - ./configs/Caddyfile:/etc/caddy/Caddyfile    # host config
→ container config
      - ./site:/usr/share/caddy                    # host site
files → container site
      - caddy_data:/data                          # named volume
→ container state
      - caddy_config:/config                      # named volume
→ container config state

volumes:
  caddy_data:
  caddy_config:
```

3. Document the volume mapping clearly

When you write docs, explain each line:

None

```
- ./configs/Caddyfile:/etc/caddy/Caddyfile      # Host config file
→ Container config
- ./site:/usr/share/caddy                        # Host HTML files →
Container site
- caddy_data:/data                              # Container state
(auto-created by Docker)
- caddy_config:/config                          # Auto-created
named volume for config state
```

5. General Rules for Docker-Compose-only Installations

- ✓ Put **all services** under `/srv/<app>`
 - ✓ Keep configs under `/srv/<app>/configs`
 - ✓ Keep site/static files under `/srv/<app>/site`
 - ✓ Use **named volumes** for container data
 - ✓ Always write a **README.md** with steps & mappings
-

◆ `/opt` vs `/srv`

`/opt`

- Stands for “**optional software**”.
- Purpose: for **add-on software packages** that are *self-contained bundles*.
- Example: if you download and install Google Chrome, MATLAB, or a vendor-provided binary → it often goes under `/opt/<package>`.
- Think of it like: “*I’m installing an external app on this machine itself.*”
- **Not intended** for per-service data, logs, configs.

👉 `/opt` is good if you're installing a *single binary/package* (like `opt/zoom`, `opt/vscode`), not managing services with configs + data.

`/srv`

- Stands for “**service data**”.
- Purpose: to hold **data and configs for network services** you run on this server.
- Example: `/srv/http` for a web server, `/srv/git` for Git repos, `/srv/mlflow` for MLflow experiments.
- Meant for **things your server serves** (apps, APIs, websites, pipelines).
- Keeps deployments clean: each service lives in `/srv/<app>` with its configs, site files, volumes, logs.

👉 `/srv` is the **right place** for Docker Compose setups because you are running services (Caddy, Airflow, MLflow, Kafka, etc.), not just installing software binaries.

◆ TL;DR

- Use `/srv` for Docker Compose deployments (Caddy, Airflow, MLflow, Kafka, RabbitMQ, etc.).
- Use `/opt` for vendor apps/binaries that aren't really “services” you are hosting, but optional software you installed.