

ML flow

1. **.env** → stores credentials and configs safely.
2. **Dockerfile.mlflow** → builds MLflow image with required libs.
3. **docker-compose.yml** → orchestrates Postgres, MinIO, bucket creation, and MLflow.

1 .env

Put this file in your project root (`~/mlflow-minio/.env`):

```
None
# Postgres
PG_USER=mlflow
PG_PASSWORD=mlflow
PG_DATABASE=mlflow
PG_PORT=5432

# MinIO
MINIO_ROOT_USER=minio
MINIO_ROOT_PASSWORD=minio123
MINIO_PORT=9000
MINIO_CONSOLE_PORT=9001
MLFLOW_BUCKET_NAME=mlflow

# MLflow
MLFLOW_PORT=5000
```

2 Dockerfile.mlflow

This goes in the root (same place as docker-compose.yml):

None

```
FROM python:3.10-slim
```

```
# Install system dependencies for psycopg2
```

```
RUN apt-get update && apt-get install -y --no-install-recommends \
```

```
    gcc libpq-dev curl && \
    rm -rf /var/lib/apt/lists/*
```

```
# Install MLflow and required libraries
```

```
RUN pip install --no-cache-dir mlflow psycopg2-binary boto3
```

```
# Default command (overridden in docker-compose.yml)
```

```
CMD ["mlflow", "server"]
```

3 docker-compose.yml

This is the full setup with **healthchecks + bucket auto-creation + networks**:

None

```
version: "3.9"
```

```
services:
```

```
  postgres:
```

```
    image: postgres:15
```

```
    container_name: mlflow_postgres
```

```
    restart: always
```

```
    environment:
```

```
      POSTGRES_USER: ${PG_USER}
```

```
      POSTGRES_PASSWORD: ${PG_PASSWORD}
```

```
      POSTGRES_DB: ${PG_DATABASE}
```

```
    volumes:
```

```
      - ./db_data:/var/lib/postgresql/data
```

```
    ports:
```

```
    - "5432:5432"
networks:
  - backend
healthcheck:
  test: ["CMD-SHELL", "pg_isready -U ${PG_USER}"]
  interval: 5s
  retries: 5

minio:
  image: minio/minio:latest
  container_name: mlflow_minio
  restart: always
  command: server /data --console-address ":9001"
  environment:
    MINIO_ROOT_USER: ${MINIO_ROOT_USER}
    MINIO_ROOT_PASSWORD: ${MINIO_ROOT_PASSWORD}
  volumes:
    - ./minio_data:/data
  ports:
    - "${MINIO_PORT}:9000"
    - "${MINIO_CONSOLE_PORT}:9001"
  networks:
    - frontend
    - backend
  healthcheck:
    test: ["CMD", "curl", "-f",
"http://localhost:9000/minio/health/live"]
    interval: 30s
    retries: 5

create_buckets:
  image: minio/mc
  container_name: mlflow_create_buckets
  depends_on:
    - minio
  entrypoint: >
```

```

    /bin/sh -c '
    sleep 10;
    mc alias set s3 http://minio:9000 ${MINIO_ROOT_USER}
    ${MINIO_ROOT_PASSWORD};
    mc mb s3/${MLFLOW_BUCKET_NAME} || echo "Bucket exists";
    mc policy set download s3/${MLFLOW_BUCKET_NAME};
    exit 0;'
networks:
  - backend

mlflow:
  build:
    context: .
    dockerfile: Dockerfile.mlflow
  container_name: mlflow_server
  restart: always
  depends_on:
    - postgres
    - minio
    - create_buckets
  ports:
    - "${MLFLOW_PORT}:5000"
  networks:
    - frontend
    - backend
  environment:
    AWS_ACCESS_KEY_ID: ${MINIO_ROOT_USER}
    AWS_SECRET_ACCESS_KEY: ${MINIO_ROOT_PASSWORD}
    MLFLOW_S3_ENDPOINT_URL: http://minio:9000
    MLFLOW_S3_IGNORE_TLS: true
  command: >
    mlflow server
    --backend-store-uri
    postgresql+psycpg2://${PG_USER}:${PG_PASSWORD}@postgres:${PG_PORT}/${PG_DATABASE}
    --default-artifact-root s3://${MLFLOW_BUCKET_NAME}

```

```
    --host 0.0.0.0
    --port 5000
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost:5000/"]
    interval: 30s
    retries: 5

volumes:
  db_data:
  minio_data:

networks:
  frontend:
  backend:
```

4 How to Start Everything

1. Save all three files (`.env`, `Dockerfile.mlflow`, `docker-compose.yml`).
2. Build and start the services:

```
Shell
docker compose up -d --build
```

3. Check logs:

```
Shell
docker compose logs -f
```

4. Verify services:

- MLflow UI → <http://localhost:5000>
- MinIO Console → <http://localhost:9001> (login: `minio/minio123`)
- Postgres running on `localhost:5432`

Great 🚀 — since you already did `docker compose up`, let's go step by step to verify **MinIO**, **MLflow**, and **Postgres**.

1 Verify MLflow Server

Check if the tracking server is running:

```
Shell
docker ps | grep mlflow_server
```

You should see a container named `mlflow_server` with port `5000->5000`.

Then open the UI in your browser:

👉 <http://localhost:5000>

If it loads, your MLflow tracking server is working ✅

2 Verify MinIO

Check if the container is running:

```
Shell
docker ps | grep mlflow_minio
```

Then open the **MinIO Console**:

👉 <http://localhost:9001>

- Username: `minio`
- Password: `minio123`

You should see a bucket called **mlflow** (created by `create_buckets` service).

👉 To check via CLI inside container:

```
Shell
docker exec -it mlflow_minio sh
```

Then:

```
Shell
mc alias set local http://localhost:9000 minio minio123
mc ls local
```

That should show the `mlflow/` bucket.

3 Verify Postgres

Check if container is running:

```
Shell
docker ps | grep mlflow_postgres
```

Then connect to Postgres:

```
Shell
docker exec -it mlflow_postgres psql -U mlflow -d mlflow
```

Inside the psql shell, run:

```
SQL
\dt
```

At first, you may see no tables. After you log some runs in MLflow, tables like **experiments**, **runs**, **metrics**, etc. will appear.

Exit with:

```
SQL
\q
```

4 Quick End-to-End Test

- This will create an experiment in **Postgres**.
- Artifacts (if logged) will go into **MinIO**.
- You'll see the run in the MLflow UI.

Testing

1. Create a folder called sample file and keep the training and testing script
2. Run this training script 1st (this will have the environment for minio) - this will create the exp and model in minio and the ml flow - verify
3. After this register the model
4. Take the run id
5. Using this run id run the testing script - it will pull the model and give the test result
6. We got all the models in minio and we got the exp and model in ml flow, and postgres data we got,
7. In this same location (sample file) we got the iris mode, test data (inside test data the have x test and y test)
8. In ml flow folder we got the db data and minio

Look like this : -


```
vikram@edgefast:~/mlflow$ ls
```

```
db_data  docker-compose.yml  Dockerfile.mlflow  minio_data  sample_file
```

```
-----
```

```
vikram@edgefast:~/mlflow/sample_file$ cat training.py
```

```
import mlflow
```

```
import mlflow.sklearn
```

```
import os
```

```
import pandas as pd
```

```
from sklearn.datasets import load_iris
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression
```

```
import joblib
```

```
# --- MLflow setup ---
```

```
mlflow.set_tracking_uri("http://192.168.1.200:5000")
```

```
mlflow.set_experiment("Iris_Classification")
```

```
# Required for artifact handling (MinIO)
```

```
os.environ["AWS_ACCESS_KEY_ID"] = "minio"
```

```
os.environ["AWS_SECRET_ACCESS_KEY"] = "minio123"
```

```
os.environ["MLFLOW_S3_ENDPOINT_URL"] = "http://192.168.1.200:9000"
```

```
def train():
```

```
    # Load dataset
```

```
iris = load_iris()

X = pd.DataFrame(iris.data, columns=iris.feature_names)

y = pd.Series(iris.target, name="target")


# Train-test split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)


# Save test data for testing script

os.makedirs("test_data", exist_ok=True)

X_test.to_csv("test_data/X_test.csv", index=False)

y_test.to_csv("test_data/y_test.csv", index=False)


# Start MLflow run

with mlflow.start_run() as run:

    print(f"Started MLflow Run ID: {run.info.run_id}")


# Train model

model = LogisticRegression(max_iter=200)

model.fit(X_train, y_train)


# Log parameters and metrics

mlflow.log_param("model_type", "LogisticRegression")
```

```

mlflow.log_param("max_iter", 200)

accuracy = model.score(X_test, y_test)

mlflow.log_metric("test_accuracy", accuracy)


# Save and log model

joblib.dump(model, "iris_model.pkl")

mlflow.log_artifact("iris_model.pkl")


# Log model using MLflow model registry format

mlflow.sklearn.log_model(model, artifact_path="model")


print(f"Model logged to MLflow with run_id={run.info.run_id}")

print(f"Test accuracy: {accuracy:.4f}")

```

```

if __name__ == "__main__":

```

```

    train()

```

```

-----

```

```

vikram@edgefast:~/mlflow/sample_file$ cat test

```

```

test_data/ testing.py

```

```

vikram@edgefast:~/mlflow/sample_file$ cat testing.py

```

```

import mlflow

```

```

import os

```

```
import pandas as pd

from sklearn.metrics import accuracy_score, classification_report


# --- MLflow setup ---

mlflow.set_tracking_uri("http://192.168.1.200:5000")


# Required for artifact handling (MinIO)

os.environ["AWS_ACCESS_KEY_ID"] = "minio"

os.environ["AWS_SECRET_ACCESS_KEY"] = "minio123"

os.environ["MLFLOW_S3_ENDPOINT_URL"] = "http://192.168.1.200:9000"


def test_latest_model():

    experiment_name = "Iris_Classification"

    client = mlflow.tracking.MlflowClient()


    # Get experiment

    experiment = client.get_experiment_by_name(experiment_name)

    if experiment is None:

        print(f"Experiment {experiment_name} not found!")

        return


    # Get latest run

    runs = client.search_runs(

        experiment_ids=[experiment.experiment_id],
```

```

        order_by=["attributes.start_time DESC"],
        max_results=1,
    )
    if not runs:
        print("No runs found!")
        return

    run = runs[0]
    run_id = run.info.run_id
    print(f"Using latest run_id: {run_id}")

    # Load model
    logged_model_uri = f"runs:{run_id}/model"
    try:
        model = mlflow.sklearn.load_model(logged_model_uri)
    except Exception as e:
        print(f"Error loading model: {e}")
        return

    # Load test data
    try:
        X_test = pd.read_csv("test_data/X_test.csv")
        y_test = pd.read_csv("test_data/y_test.csv")
    except FileNotFoundError:

```

```

    print("Error: test data not found. Run training first!")

    return

# Evaluate

y_pred = model.predict(X_test)

acc = accuracy_score(y_test, y_pred)

report = classification_report(y_test, y_pred)


print("\n--- Test Results ---")

print(f"Test Accuracy: {acc:.4f}")

print("\nClassification Report:")

print(report)

print("-----")


if __name__ == "__main__":

    test_latest_model()

vikram@edgefast:~/mlflow/sample_file$

-----

```

Afer running the training you will see the models are in mino and the ml flow after that you run the testing script udin the run id thats it - make sure the postgress is accessible in the pgadmin software and also it has that datas are there

Also check inside the container the datas are available test need to add hera

1. Find your Postgres container

To confirm:

Shell

```
docker ps --format "table {{.Names}}\t{{.Image}}\t{{.Status}}"
```

Look for the container running `postgres`.

From your `docker ps` output, your **Postgres container name is**:

None

`mlflow_postgres`

Use: (not use)

Shell

```
sudo docker exec -it mlflow_postgres psql -U $PG_USER -d  
$PG_DATABASE
```

If your `.env` has:

None

`PG_USER=mlflow`

`PG_DATABASE=mlflow`

then run: (use this)

Shell

```
sudo docker exec -it mlflow_postgres psql -U mlflow -d mlflow
```

Once inside psql

- List tables:

```
SQL  
\dt
```

- Check experiments:

```
SQL  
SELECT * FROM experiments;
```

- Check runs:

```
SQL  
SELECT run_uuid, experiment_id, status, start_time FROM runs  
LIMIT 10;
```

- Check schema of `runs` table:

```
SQL  
\d runs
```

```
sudo mkdir mlflow
699 ls
700 cd mlflow/
701 ls
702 sudo nano .env
703 sudo nano Dockerfile.mlflow
704 sudo nano docker-compose.yml
705 ls
706 ll
707 docker compose up -d --build
708 sudo docker compose up -d --build
709 docker compose logs -f
710 sudo docker compose logs -f
711 sudo docker exec -it mlflow_minio sh
712 docker ps | grep mlflow_postgres
713 sudo docker ps | grep mlflow_postgres
714 docker exec -it mlflow_postgres psql -U mlflow -d mlflow
715 sudo docker exec -it mlflow_postgres psql -U mlflow -d mlflow
716 ls
717 cat docker-compose.yml
718 ls
719 ll
720 cd minio_data/
721 ls
722 ls mlflow/
723 cd mlflow/
724 ls
725 cd ..
726 ls
727 cd db_data/
728 sudo ls db_data/
729 ls
730 ll
731 ls
732 sudo mkdir sample_file
733 ls
734 cd sample_file/
735 ls
736 sudo nano training.py
737 python training.py
738 sudo chown -R vikram:vikram ~/mlflow/sample_file
739 python training.py
740 ls
741 cd test_data/
```

```
742 ls
743 cd ..
744 sudo docker exec -it mlflow_postgres psql -U mlflow -d mlflow
745 ls
746 sudo nano testing.py
747 ls
748 python3 testing.py
749 history
vikram@edgefast:~/mlflow/sample_file$
```